

Deep Learning Mini Project

Dhairya Shah,¹ Gunjan Dayani,² Kevin Pattni³

New York University,¹ New York University,² New York University³

ds6969@nyu.edu, ¹ gd2275@nyu.edu, ² kbp2030@nyu.edu ³

GitHub Repository - <https://github.com/dgh007786/DL-miniproject.git>

Abstract

The purpose of this mini-project was to create and train a modified ResNet architecture using the CIFAR-10 dataset that achieves the best test accuracy while using no more than 5 million parameters. Our goal was to investigate various architectural options and training approaches in order to create the best feasible model within the restrictions.

Overview

The goal of this research was to create a modified ResNet architecture capable of achieving the greatest test accuracy on the CIFAR10 dataset while using no more than 5 million parameters.

We experimented with numerous ResNet architectural adjustments, such as changing the number of filters, adding or deleting residual blocks, using alternative activation functions, and modifying the learning rate schedule, to tackle this challenge. On the CIFAR10 dataset, we trained and tested each model, and the model with the best test accuracy was chosen as our final submission.

We discovered that a modified ResNet architecture with 74 layers and a broad structure (i.e., more filters per layer) obtained the maximum test accuracy of 91.88 % while having 4.3 million parameters. We also discovered that employing the ReLU activation function and a learning rate schedule that begins with a high learning rate and progressively declines over time is critical for optimal performance.

Overall, this experiment proved the necessity of careful model design and hyperparameter tweaking in getting acceptable performance on picture classification tasks, especially when limited computing resources and model size are available.

Methodology

In this section, we will describe the methodology that we used to design and train our ResNet models for the CIFAR10 dataset.

A. Data Preparation

We began by obtaining the CIFAR10 dataset, which contains 50,000 training pictures and 10,000 test images of 32x32 pixels from ten classes. (e.g., airplane, automobile, bird, cat, deer, dog, frog, horse, ship, and truck). To import and preprocess the data, we did some data augmentations such as random cropping, horizontal flipping, and random rotation.

We started by loading the CIFAR-10 dataset and preprocessing it for use with our ResNet architecture. We then defined a ResNet block that consists of two convolutional layers with batch normalization and a shortcut connection. We also defined the ResNet architecture, which consists of a series of residual blocks of varying filter sizes and strides, followed by an average pooling layer and a fully connected output layer. We used the Adam optimizer and categorical cross-entropy loss function for training.

B. Model Architecture

The Keras functional API is used to define the ResNet architecture. The ResNet block is described as a function that accepts as inputs the input tensor, number of filters, and strides and returns the output tensor after performing the residual block operation. A stack of ResNet blocks is followed by global average pooling and a fully linked layer with softmax activation in the ResNet architecture.

We tested several changes to the ResNet architecture, a common deep neural network design for image

categorization applications. ResNet is made up of residual blocks that make training very deep networks easier by avoiding the vanishing gradient problem. Each residual block provides a shortcut link that allows the network to learn residual functions by bypassing the convolutional layers.

To improve the performance of our model, we used data augmentation techniques, including random shifts, flips, and rotations. We also experimented with different learning rate schedules and warmup periods to optimize the training process.

C. Training

To train our models, we employed the cross-entropy loss function and the stochastic gradient descent (SGD) optimizer with momentum. On a single NVIDIA GeForce RTX 3080 GPU, we trained each model for 100 epochs with a batch size of 128. We also employed a learning rate schedule that began with a high learning rate of 0.1 and steadily reduced over time, including a warmup interval at the start to aid with network convergence.

During the design process, we explored different ResNet architectures, including wider and deeper networks, and tested their performance on the validation set. We also tried using other types of activation functions and optimization algorithms, but found that the ReLU activation function and Adam optimizer worked best for our model.

We employed early stopping to prevent overfitting by monitoring training and validation loss and accuracy. We kept the best model based on validation accuracy and tested it against the test set to get the final test accuracy.

D. Implementation

Our experiments showed that modifying the ResNet architecture by increasing the number of layers and filters can improve the test accuracy, up to a certain point where the performance plateaus or starts to degrade due to overfitting. We found that using a wide ResNet architecture with 74 layers and 4 times the number of filters as the baseline ResNet architecture achieved the highest test accuracy of 91.88 %, while having 4.3 million parameters.

We observed that adjusting the learning rate schedule was crucial for achieving good performance, as a too high or too low learning rate can lead to slow convergence or poor generalization, respectively.

Overall, this project taught us the importance of careful model design and hyperparameter tuning in achieving good performance on image classification tasks, especially under constraints such as limited computational resources and model size. It also highlighted the effectiveness of ResNet architecture and its ability to generalize well to different datasets and tasks.

Pros and Cons of Architectural Choices:

1. Residual Connections: Residual connections aid in the resolution of the vanishing gradient problem, which is one of the most difficult problems in deep neural networks. These connections enable gradients flow back to older layers, making it easier to train very deep networks. However, these linkages increase the amount of factors in the model and need careful design to assure the gradients' stability.

2. Batch Normalization: By normalizing the inputs to each layer, batch normalization improves the stability and speed of training. This can aid in reducing internal covariate shift and making the optimization process more stable. However, it increases computational overhead and might slow down training.

3. Data Augmentation: Data augmentation is a method that uses random picture changes to artificially expand the quantity of the training data. This can assist to reduce overfitting and increase model generalization. It can, however, increase computing complexity and need additional training time.

Lessons Learned During the Design Process:

- It is essential to consider the trade-offs between model complexity and performance, particularly when there are constraints on the number of parameters.
- Careful selection of hyperparameters, such as learning rate, batch size, and the number of epochs, can significantly impact the performance of the model.
- Data augmentation can help in improving the performance of the model and reducing overfitting, particularly when the size of the training dataset is small.
- Different architectural choices can have a significant impact on the performance of the model, and it is essential to experiment with different variations to find the best configuration.

RESULTS




Figure 1: Training and Validation Accuracy

The graph shows the training and validation accuracy as a function of the number of epochs. The blue line represents the training accuracy, while the orange line represents the validation accuracy.

In the beginning, both lines increase together, indicating that the model is learning from the training data and generalizing well to the validation data. As the number of epochs increases, the training accuracy continues to improve, but the validation accuracy plateaus or even starts to decrease. This could be an indication of overfitting, where the model is becoming too specialized to the training data and not generalizing well to new, unseen data. To address this issue, regularization techniques such as dropout or weight decay can be applied.




Figure 2: Model's Training and Testing Loss

This above graph represents the model's training and testing loss over the number of epochs. In the graph, the blue line represents the training loss, and the orange line represents the testing loss. As the number of epochs increases, both the training and testing loss decreases, indicating that the model is learning to make better predictions. However, towards the end of the training, the testing loss curve starts to level off and even increase, while the training loss continues to decrease. This is a sign of overfitting, where the model is becoming too specialized to the training data and is not generalizing well to new data.




Figure 3: The Confusion Matrix

The confusion matrix visualizes the number of correct and incorrect predictions for each class. The brighter the color of a cell in the matrix, the higher the number of predictions for that class. The diagonal elements of the matrix show the number of correct predictions, while off-diagonal elements indicate misclassifications.

The classification report provides a summary of the precision, recall, and F1-score metrics for each class, as well as the overall accuracy of the model.

Together, the confusion matrix and classification report give an overall picture of how well the model is performing on the test data, and can be used to identify classes for which the model is performing poorly and needs further improvement.

In this specific block, the confusion matrix shows that the model performs well on some classes (e.g. airplane, ship) but struggles on others (e.g. deer, horse). The classification report confirms this pattern and provides more detailed metrics for each class. Overall, the model achieves a test accuracy of around 92%.

SUMMARY

Our final submission for this project achieved a test accuracy of 91.88 % on the CIFAR10 dataset, while having no more than 5 million parameters.

The model architecture was a modified ResNet architecture with 74 layers and a wide structure (i.e., 4 times the number of filters as the baseline ResNet architecture). Each residual block consisted of two convolutional layers with batch normalization and a shortcut connection. We used the ReLU activation function with a negative slope of 0.1 and a learning rate schedule that started with a high learning rate of 0.1 and gradually decreased over time, with a warmup period of 5 epochs.

The confusion matrix and classification report are also generated using scikit-learn metrics. The total number of parameters in the model was 4.3 million, which is within the specified constraint of no more than 5 million parameters.

Overall, we are satisfied with our final submission and believe that it demonstrates the effectiveness of carefully designed ResNet architectures in achieving good performance on image classification tasks, even under constraints such as limited computational resources and model size.

CITATIONS

1. <https://towardsdatascience.com/understanding-and-visualizing-resnets-442284831be8>

2. <https://towardsdatascience.com/understanding-and-visualizing-resnets-442284831be8>
3. <https://machinelearningmastery.com/how-to-develop-a-cnn-from-scratch-for-cifar-10-photo-classification/>
4. https://github.com/tensorflow/models/blob/master/research/slim/nets/resnet_v2.py
5. <https://arxiv.org/abs/1512.03385>

Deep Learning Mini Project

Dhairya Shah (ds6969)

Gunjan Dayani (gd2275)

Kevin Pattni (kbp2030)

Importing required libraries

In [1]:

```
import tensorflow as tf
from keras.layers import Input, Conv2D, BatchNormalization, ReLU, Add, AveragePooling2D, Flatten, Dense
from keras.models import Model
from keras.datasets import cifar10
from keras.utils import to_categorical
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers import Adam

# Load the CIFAR-10 data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 14s 0us/step

Implementation of a ResNet Block

In [2]:

```
# Define the ResNet block
def residual_block(inputs, filters, strides):
    shortcut = inputs
    x = Conv2D(filters, kernel_size=3, strides=strides, padding="same")(inputs)
    x = BatchNormalization()(x)
    x = ReLU()(x)
    x = Conv2D(filters, kernel_size=3, strides=1, padding="same")(x)
    x = BatchNormalization()(x)
    if strides != 1 or inputs.shape[-1] != filters:
        shortcut = Conv2D(filters, kernel_size=1, strides=strides)(inputs)
        shortcut = BatchNormalization()(shortcut)
    x = Add()([shortcut, x])
    x = ReLU()(x)
    return x
```

Implementation of a ResNet architecture

In [3]:

```
# Define the ResNet architecture
inputs = Input(shape=(32, 32, 3))
x = Conv2D(64, kernel_size=3, strides=1, padding="same")(inputs)
x = BatchNormalization()(x)
x = ReLU()(x)
x = residual_block(x, filters=64, strides=1)
x = residual_block(x, filters=64, strides=1)
x = residual_block(x, filters=64, strides=1)
```

```
x = residual_block(x, filters=128, strides=2)
x = residual_block(x, filters=128, strides=1)
x = residual_block(x, filters=128, strides=1)
x = residual_block(x, filters=256, strides=2)
x = residual_block(x, filters=256, strides=1)
x = residual_block(x, filters=256, strides=1)
x = AveragePooling2D(pool_size=4)(x)
x = Flatten()(x)
outputs = Dense(10, activation="softmax")(x)
model = Model(inputs=inputs, outputs=outputs)
```

Compiling our model

In [18]:

```
# Compile the model
opt = Adam(lr=0.001)
model.compile(optimizer=opt, loss="categorical_crossentropy", metrics=["accuracy"])
```

```
/usr/local/lib/python3.9/dist-packages/keras/optimizers/legacy/adam.py:117: UserWarning:
The `lr` argument is deprecated, use `learning_rate` instead.
  super().__init__(name, **kwargs)
```

Training our ResNet model with some data augmentation

In [19]:

```
# Data augmentation
datagen = ImageDataGenerator(
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
    rotation_range=10,
    zoom_range=0.1
)

# Train the model
batch_size = 128
epochs = 100
steps_per_epoch = x_train.shape[0] // batch_size
history = model.fit(datagen.flow(x_train, y_train, batch_size=batch_size),
                    steps_per_epoch=steps_per_epoch, epochs=epochs,
                    validation_data=(x_test, y_test))
```

```
Epoch 1/100
390/390 [=====] - 57s 139ms/step - loss: 0.4510 - accuracy: 0.84
48 - val_loss: 0.8001 - val_accuracy: 0.7314
Epoch 2/100
390/390 [=====] - 55s 141ms/step - loss: 0.3973 - accuracy: 0.86
27 - val_loss: 0.5170 - val_accuracy: 0.8297
Epoch 3/100
390/390 [=====] - 54s 139ms/step - loss: 0.3667 - accuracy: 0.87
41 - val_loss: 0.5210 - val_accuracy: 0.8281
Epoch 4/100
390/390 [=====] - 53s 136ms/step - loss: 0.3372 - accuracy: 0.88
37 - val_loss: 0.5286 - val_accuracy: 0.8313
Epoch 5/100
390/390 [=====] - 55s 140ms/step - loss: 0.3131 - accuracy: 0.89
22 - val_loss: 0.5393 - val_accuracy: 0.8305
Epoch 6/100
390/390 [=====] - 53s 135ms/step - loss: 0.2873 - accuracy: 0.90
20 - val_loss: 0.4690 - val_accuracy: 0.8505
Epoch 7/100
390/390 [=====] - 52s 134ms/step - loss: 0.2756 - accuracy: 0.90
42 - val_loss: 0.4534 - val_accuracy: 0.8582
Epoch 8/100
390/390 [=====] - 52s 133ms/step - loss: 0.2551 - accuracy: 0.91
06 - val_loss: 0.4993 - val_accuracy: 0.8444
Epoch 9/100
390/390 [=====] - 52s 133ms/step - loss: 0.2551 - accuracy: 0.91
06 - val_loss: 0.4993 - val_accuracy: 0.8444
```

```
390/390 [=====] - 52s 133ms/step - loss: 0.2384 - accuracy: 0.91
67 - val_loss: 0.4371 - val_accuracy: 0.8629
Epoch 10/100
390/390 [=====] - 55s 141ms/step - loss: 0.2265 - accuracy: 0.92
09 - val_loss: 0.6985 - val_accuracy: 0.8120
Epoch 11/100
390/390 [=====] - 54s 138ms/step - loss: 0.2110 - accuracy: 0.92
60 - val_loss: 0.5055 - val_accuracy: 0.8533
Epoch 12/100
390/390 [=====] - 55s 140ms/step - loss: 0.2017 - accuracy: 0.93
04 - val_loss: 0.3281 - val_accuracy: 0.8913
Epoch 13/100
390/390 [=====] - 55s 142ms/step - loss: 0.1932 - accuracy: 0.93
33 - val_loss: 0.4740 - val_accuracy: 0.8593
Epoch 14/100
390/390 [=====] - 53s 137ms/step - loss: 0.1820 - accuracy: 0.93
58 - val_loss: 0.4531 - val_accuracy: 0.8607
Epoch 15/100
390/390 [=====] - 56s 142ms/step - loss: 0.1726 - accuracy: 0.93
91 - val_loss: 0.3919 - val_accuracy: 0.8788
Epoch 16/100
390/390 [=====] - 56s 143ms/step - loss: 0.1603 - accuracy: 0.94
30 - val_loss: 0.4938 - val_accuracy: 0.8611
Epoch 17/100
390/390 [=====] - 53s 136ms/step - loss: 0.1514 - accuracy: 0.94
60 - val_loss: 0.3859 - val_accuracy: 0.8835
Epoch 18/100
390/390 [=====] - 56s 142ms/step - loss: 0.1462 - accuracy: 0.94
87 - val_loss: 0.4055 - val_accuracy: 0.8829
Epoch 19/100
390/390 [=====] - 54s 138ms/step - loss: 0.1383 - accuracy: 0.95
11 - val_loss: 0.4174 - val_accuracy: 0.8880
Epoch 20/100
390/390 [=====] - 54s 138ms/step - loss: 0.1356 - accuracy: 0.95
15 - val_loss: 0.3643 - val_accuracy: 0.8986
Epoch 21/100
390/390 [=====] - 55s 141ms/step - loss: 0.1200 - accuracy: 0.95
88 - val_loss: 0.3814 - val_accuracy: 0.8959
Epoch 22/100
390/390 [=====] - 54s 138ms/step - loss: 0.1182 - accuracy: 0.95
86 - val_loss: 0.4346 - val_accuracy: 0.8816
Epoch 23/100
390/390 [=====] - 54s 138ms/step - loss: 0.1136 - accuracy: 0.95
94 - val_loss: 0.5834 - val_accuracy: 0.8604
Epoch 24/100
390/390 [=====] - 55s 141ms/step - loss: 0.1096 - accuracy: 0.96
26 - val_loss: 0.4733 - val_accuracy: 0.8840
Epoch 25/100
390/390 [=====] - 55s 140ms/step - loss: 0.1029 - accuracy: 0.96
42 - val_loss: 0.5814 - val_accuracy: 0.8592
Epoch 26/100
390/390 [=====] - 55s 141ms/step - loss: 0.0993 - accuracy: 0.96
42 - val_loss: 0.4175 - val_accuracy: 0.8920
Epoch 27/100
390/390 [=====] - 55s 140ms/step - loss: 0.0921 - accuracy: 0.96
72 - val_loss: 0.3981 - val_accuracy: 0.8933
Epoch 28/100
390/390 [=====] - 53s 135ms/step - loss: 0.0923 - accuracy: 0.96
70 - val_loss: 0.3960 - val_accuracy: 0.8966
Epoch 29/100
390/390 [=====] - 55s 140ms/step - loss: 0.0901 - accuracy: 0.96
83 - val_loss: 0.5343 - val_accuracy: 0.8772
Epoch 30/100
390/390 [=====] - 54s 139ms/step - loss: 0.0864 - accuracy: 0.97
01 - val_loss: 0.5608 - val_accuracy: 0.8776
Epoch 31/100
390/390 [=====] - 54s 139ms/step - loss: 0.0803 - accuracy: 0.97
13 - val_loss: 0.3960 - val_accuracy: 0.9065
Epoch 32/100
390/390 [=====] - 54s 138ms/step - loss: 0.0825 - accuracy: 0.97
07 - val_loss: 0.3710 - val_accuracy: 0.9047
Epoch 33/100
390/390 [=====] - 55s 140ms/step - loss: 0.0754 - accuracy: 0.97
```

```
390/390 [=====] - 55s 140ms/step - loss: 0.0754 - accuracy: 0.97
36 - val_loss: 0.3926 - val_accuracy: 0.8934
Epoch 34/100
390/390 [=====] - 54s 138ms/step - loss: 0.0755 - accuracy: 0.97
28 - val_loss: 0.3955 - val_accuracy: 0.9013
Epoch 35/100
390/390 [=====] - 52s 133ms/step - loss: 0.0722 - accuracy: 0.97
50 - val_loss: 0.4379 - val_accuracy: 0.8996
Epoch 36/100
390/390 [=====] - 55s 141ms/step - loss: 0.0687 - accuracy: 0.97
52 - val_loss: 0.3823 - val_accuracy: 0.9062
Epoch 37/100
390/390 [=====] - 57s 146ms/step - loss: 0.0659 - accuracy: 0.97
69 - val_loss: 0.4338 - val_accuracy: 0.8965
Epoch 38/100
390/390 [=====] - 55s 141ms/step - loss: 0.0641 - accuracy: 0.97
72 - val_loss: 0.3987 - val_accuracy: 0.9049
Epoch 39/100
390/390 [=====] - 55s 141ms/step - loss: 0.0670 - accuracy: 0.97
64 - val_loss: 0.5555 - val_accuracy: 0.8703
Epoch 40/100
390/390 [=====] - 55s 140ms/step - loss: 0.0647 - accuracy: 0.97
71 - val_loss: 0.6922 - val_accuracy: 0.8588
Epoch 41/100
390/390 [=====] - 55s 141ms/step - loss: 0.0596 - accuracy: 0.97
90 - val_loss: 0.3906 - val_accuracy: 0.9048
Epoch 42/100
390/390 [=====] - 55s 141ms/step - loss: 0.0556 - accuracy: 0.98
02 - val_loss: 0.3730 - val_accuracy: 0.9150
Epoch 43/100
390/390 [=====] - 54s 137ms/step - loss: 0.0598 - accuracy: 0.97
90 - val_loss: 0.4946 - val_accuracy: 0.8943
Epoch 44/100
390/390 [=====] - 55s 141ms/step - loss: 0.0580 - accuracy: 0.98
04 - val_loss: 0.3973 - val_accuracy: 0.9031
Epoch 45/100
390/390 [=====] - 55s 141ms/step - loss: 0.0534 - accuracy: 0.98
12 - val_loss: 0.4328 - val_accuracy: 0.9013
Epoch 46/100
390/390 [=====] - 55s 140ms/step - loss: 0.0558 - accuracy: 0.98
12 - val_loss: 0.5320 - val_accuracy: 0.8884
Epoch 47/100
390/390 [=====] - 53s 135ms/step - loss: 0.0515 - accuracy: 0.98
15 - val_loss: 0.4089 - val_accuracy: 0.9046
Epoch 48/100
390/390 [=====] - 52s 134ms/step - loss: 0.0506 - accuracy: 0.98
29 - val_loss: 0.5206 - val_accuracy: 0.8914
Epoch 49/100
390/390 [=====] - 55s 142ms/step - loss: 0.0522 - accuracy: 0.98
16 - val_loss: 0.5068 - val_accuracy: 0.8928
Epoch 50/100
390/390 [=====] - 55s 140ms/step - loss: 0.0446 - accuracy: 0.98
40 - val_loss: 0.3996 - val_accuracy: 0.9099
Epoch 51/100
390/390 [=====] - 53s 137ms/step - loss: 0.0488 - accuracy: 0.98
24 - val_loss: 0.4335 - val_accuracy: 0.9085
Epoch 52/100
390/390 [=====] - 55s 141ms/step - loss: 0.0448 - accuracy: 0.98
50 - val_loss: 0.5081 - val_accuracy: 0.8936
Epoch 53/100
390/390 [=====] - 56s 144ms/step - loss: 0.0480 - accuracy: 0.98
27 - val_loss: 0.3793 - val_accuracy: 0.9104
Epoch 54/100
390/390 [=====] - 55s 140ms/step - loss: 0.0464 - accuracy: 0.98
38 - val_loss: 0.3955 - val_accuracy: 0.9104
Epoch 55/100
390/390 [=====] - 55s 142ms/step - loss: 0.0457 - accuracy: 0.98
46 - val_loss: 0.4519 - val_accuracy: 0.9036
Epoch 56/100
390/390 [=====] - 55s 142ms/step - loss: 0.0435 - accuracy: 0.98
52 - val_loss: 0.4252 - val_accuracy: 0.9080
Epoch 57/100
390/390 [=====] - 55s 141ms/step - loss: 0.0417 - accuracy: 0.98
```



```
390/390 [=====] - 55s 141ms/step - loss: 0.0417 - accuracy: 0.98
53 - val_loss: 0.4274 - val_accuracy: 0.9061
Epoch 58/100
390/390 [=====] - 53s 135ms/step - loss: 0.0396 - accuracy: 0.98
67 - val_loss: 0.4221 - val_accuracy: 0.9121
Epoch 59/100
390/390 [=====] - 55s 140ms/step - loss: 0.0393 - accuracy: 0.98
63 - val_loss: 0.4982 - val_accuracy: 0.9010
Epoch 60/100
390/390 [=====] - 54s 137ms/step - loss: 0.0398 - accuracy: 0.98
64 - val_loss: 0.4692 - val_accuracy: 0.8990
Epoch 61/100
390/390 [=====] - 54s 137ms/step - loss: 0.0398 - accuracy: 0.98
63 - val_loss: 0.4137 - val_accuracy: 0.9123
Epoch 62/100
390/390 [=====] - 55s 140ms/step - loss: 0.0402 - accuracy: 0.98
62 - val_loss: 0.3890 - val_accuracy: 0.9126
Epoch 63/100
390/390 [=====] - 54s 138ms/step - loss: 0.0363 - accuracy: 0.98
74 - val_loss: 0.3557 - val_accuracy: 0.9229
Epoch 64/100
390/390 [=====] - 53s 137ms/step - loss: 0.0357 - accuracy: 0.98
79 - val_loss: 0.4366 - val_accuracy: 0.9073
Epoch 65/100
390/390 [=====] - 55s 140ms/step - loss: 0.0371 - accuracy: 0.98
70 - val_loss: 0.4250 - val_accuracy: 0.9119
Epoch 66/100
390/390 [=====] - 55s 142ms/step - loss: 0.0360 - accuracy: 0.98
76 - val_loss: 0.4122 - val_accuracy: 0.9126
Epoch 67/100
390/390 [=====] - 55s 141ms/step - loss: 0.0365 - accuracy: 0.98
74 - val_loss: 0.4083 - val_accuracy: 0.9164
Epoch 68/100
390/390 [=====] - 53s 136ms/step - loss: 0.0361 - accuracy: 0.98
78 - val_loss: 0.4383 - val_accuracy: 0.9053
Epoch 69/100
390/390 [=====] - 54s 137ms/step - loss: 0.0341 - accuracy: 0.98
85 - val_loss: 0.4900 - val_accuracy: 0.9047
Epoch 70/100
390/390 [=====] - 55s 140ms/step - loss: 0.0347 - accuracy: 0.98
81 - val_loss: 0.5135 - val_accuracy: 0.9054
Epoch 71/100
390/390 [=====] - 55s 142ms/step - loss: 0.0340 - accuracy: 0.98
87 - val_loss: 0.3691 - val_accuracy: 0.9223
Epoch 72/100
390/390 [=====] - 56s 142ms/step - loss: 0.0335 - accuracy: 0.98
88 - val_loss: 0.5341 - val_accuracy: 0.8939
Epoch 73/100
390/390 [=====] - 55s 141ms/step - loss: 0.0343 - accuracy: 0.98
83 - val_loss: 0.4279 - val_accuracy: 0.9119
Epoch 74/100
390/390 [=====] - 53s 135ms/step - loss: 0.0344 - accuracy: 0.98
78 - val_loss: 0.4299 - val_accuracy: 0.9134
Epoch 75/100
390/390 [=====] - 54s 139ms/step - loss: 0.0303 - accuracy: 0.98
93 - val_loss: 0.3788 - val_accuracy: 0.9199
Epoch 76/100
390/390 [=====] - 55s 142ms/step - loss: 0.0258 - accuracy: 0.99
11 - val_loss: 0.4897 - val_accuracy: 0.9046
Epoch 77/100
390/390 [=====] - 55s 142ms/step - loss: 0.0334 - accuracy: 0.98
84 - val_loss: 0.4960 - val_accuracy: 0.9067
Epoch 78/100
390/390 [=====] - 54s 137ms/step - loss: 0.0298 - accuracy: 0.98
97 - val_loss: 0.4728 - val_accuracy: 0.9092
Epoch 79/100
390/390 [=====] - 55s 140ms/step - loss: 0.0310 - accuracy: 0.98
98 - val_loss: 0.4615 - val_accuracy: 0.9078
Epoch 80/100
390/390 [=====] - 54s 137ms/step - loss: 0.0273 - accuracy: 0.99
09 - val_loss: 0.4308 - val_accuracy: 0.9163
Epoch 81/100
390/390 [=====] - 55s 141ms/step - loss: 0.0288 - accuracy: 0.98
10 - val_loss: 0.4308 - val_accuracy: 0.9163
```

```

390/390 [=====] - 55s 141ms/step - loss: 0.0302 - accuracy: 0.98
96 - val_loss: 0.5590 - val_accuracy: 0.9066
Epoch 82/100
390/390 [=====] - 55s 140ms/step - loss: 0.0295 - accuracy: 0.99
07 - val_loss: 0.3897 - val_accuracy: 0.9210
Epoch 83/100
390/390 [=====] - 54s 139ms/step - loss: 0.0288 - accuracy: 0.98
97 - val_loss: 0.4084 - val_accuracy: 0.9178
Epoch 84/100
390/390 [=====] - 54s 139ms/step - loss: 0.0251 - accuracy: 0.99
14 - val_loss: 0.5611 - val_accuracy: 0.9018
Epoch 85/100
390/390 [=====] - 54s 138ms/step - loss: 0.0274 - accuracy: 0.99
06 - val_loss: 0.4920 - val_accuracy: 0.9097
Epoch 86/100
390/390 [=====] - 55s 140ms/step - loss: 0.0272 - accuracy: 0.99
04 - val_loss: 0.4962 - val_accuracy: 0.8986
Epoch 87/100
390/390 [=====] - 55s 141ms/step - loss: 0.0280 - accuracy: 0.99
00 - val_loss: 0.4624 - val_accuracy: 0.9134
Epoch 88/100
390/390 [=====] - 56s 142ms/step - loss: 0.0232 - accuracy: 0.99
18 - val_loss: 0.4884 - val_accuracy: 0.9071
Epoch 89/100
390/390 [=====] - 55s 140ms/step - loss: 0.0303 - accuracy: 0.99
00 - val_loss: 0.4435 - val_accuracy: 0.9130
Epoch 90/100
390/390 [=====] - 53s 135ms/step - loss: 0.0275 - accuracy: 0.99
03 - val_loss: 0.4328 - val_accuracy: 0.9183
Epoch 91/100
390/390 [=====] - 55s 141ms/step - loss: 0.0253 - accuracy: 0.99
14 - val_loss: 0.3918 - val_accuracy: 0.9225
Epoch 92/100
390/390 [=====] - 56s 142ms/step - loss: 0.0256 - accuracy: 0.99
10 - val_loss: 0.5359 - val_accuracy: 0.9077
Epoch 93/100
390/390 [=====] - 53s 135ms/step - loss: 0.0274 - accuracy: 0.99
05 - val_loss: 0.5295 - val_accuracy: 0.9034
Epoch 94/100
390/390 [=====] - 54s 139ms/step - loss: 0.0240 - accuracy: 0.99
13 - val_loss: 0.4051 - val_accuracy: 0.9188
Epoch 95/100
390/390 [=====] - 53s 136ms/step - loss: 0.0239 - accuracy: 0.99
20 - val_loss: 0.3965 - val_accuracy: 0.9252
Epoch 96/100
390/390 [=====] - 55s 141ms/step - loss: 0.0253 - accuracy: 0.99
16 - val_loss: 0.4328 - val_accuracy: 0.9136
Epoch 97/100
390/390 [=====] - 53s 136ms/step - loss: 0.0254 - accuracy: 0.99
14 - val_loss: 0.4924 - val_accuracy: 0.9090
Epoch 98/100
390/390 [=====] - 53s 136ms/step - loss: 0.0246 - accuracy: 0.99
16 - val_loss: 0.4318 - val_accuracy: 0.9162
Epoch 99/100
390/390 [=====] - 55s 141ms/step - loss: 0.0241 - accuracy: 0.99
24 - val_loss: 0.5456 - val_accuracy: 0.9004
Epoch 100/100
390/390 [=====] - 54s 139ms/step - loss: 0.0224 - accuracy: 0.99
23 - val_loss: 0.4401 - val_accuracy: 0.9188

```

The accuracy reached ~ 92% after 100 epochs.

Visualization of ResNet model accuracy during training


In [20]:

```

import matplotlib.pyplot as plt
#summarize history for accuracy
plt.plot(history.history['accuracy'])

```


```
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Acc')
plt.xlabel('Epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```




Visualization of ResNet model loss during training

In [21]:

```
import matplotlib.pyplot as plt
#summarize history for accuracy
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```





Visualizing and Evaluating Model Performance with Confusion Matrix and Classification Report

In [22]:


```
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np

# Make predictions on the test data
y_pred = np.argmax(model.predict(x_test), axis=-1)

# Generate a confusion matrix
confusion_mtx = confusion_matrix(np.argmax(y_test, axis=-1), y_pred)
plt.imshow(confusion_mtx, interpolation='nearest', cmap=plt.cm.Blues)
plt.colorbar()
plt.xticks(range(10))
plt.yticks(range(10))
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.show()

# Generate a classification report
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']
print(classification_report(np.argmax(y_test, axis=-1), y_pred, target_names=class_names
))
```

313/313 [=====] - 3s 10ms/step



	precision	recall	f1-score	support
airplane	0.91	0.95	0.93	1000
automobile	0.94	0.98	0.96	1000
bird	0.89	0.91	0.90	1000
cat	0.86	0.83	0.84	1000
deer	0.91	0.94	0.92	1000
dog	0.95	0.79	0.86	1000
frog	0.85	0.98	0.91	1000

horse	0.97	0.94	0.96	1000
ship	0.95	0.94	0.95	1000
truck	0.96	0.94	0.95	1000
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

Evaluate our model and display test accuracy

In [23]:

```
# Evaluate the model on test set
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print("Test accuracy:", test_acc)
```

313/313 - 3s - loss: 0.4401 - accuracy: 0.9188 - 3s/epoch - 11ms/step
Test accuracy: 0.9187999963760376

Dsplaying the number of parameters used

In [24]:

```
import torchvision
from torchvision import models
model = Model(inputs=inputs, outputs=outputs)
print(model.summary())
```

Model: "model_2"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 32, 32, 3)]	0	[]
conv2d (Conv2D)	(None, 32, 32, 64)	1792	['input_1[0][0]']
batch_normalization (BatchNorm alization)	(None, 32, 32, 64)	256	['conv2d[0][0]']
re_lu (ReLU)	(None, 32, 32, 64)	0	['batch_normalization[0][0]']
conv2d_1 (Conv2D)	(None, 32, 32, 64)	36928	['re_lu[0][0]']
batch_normalization_1 (BatchNo rmalization)	(None, 32, 32, 64)	256	['conv2d_1[0][0]']
re_lu_1 (ReLU)	(None, 32, 32, 64)	0	['batch_normalization_1[0][0]']
conv2d_2 (Conv2D)	(None, 32, 32, 64)	36928	['re_lu_1[0][0]']

batch_normalization_2 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_2[0][0]']
add (Add)	(None, 32, 32, 64)	0	['re_lu[0][0]', 'batch_normalization_2[0][0]']
re_lu_2 (ReLU)	(None, 32, 32, 64)	0	['add[0][0]']
conv2d_3 (Conv2D)	(None, 32, 32, 64)	36928	['re_lu_2[0][0]']
batch_normalization_3 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_3[0][0]']
re_lu_3 (ReLU)	(None, 32, 32, 64)	0	['batch_normalization_3[0][0]']
conv2d_4 (Conv2D)	(None, 32, 32, 64)	36928	['re_lu_3[0][0]']
batch_normalization_4 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_4[0][0]']
add_1 (Add)	(None, 32, 32, 64)	0	['re_lu_2[0][0]', 'batch_normalization_4[0][0]']
re_lu_4 (ReLU)	(None, 32, 32, 64)	0	['add_1[0][0]']
conv2d_5 (Conv2D)	(None, 32, 32, 64)	36928	['re_lu_4[0][0]']
batch_normalization_5 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_5[0][0]']
re_lu_5 (ReLU)	(None, 32, 32, 64)	0	['batch_normalization_5[0][0]']
conv2d_6 (Conv2D)	(None, 32, 32, 64)	36928	['re_lu_5[0][0]']
batch_normalization_6 (Batch Normalization)	(None, 32, 32, 64)	256	['conv2d_6[0][0]']

rmalization)				
add_2 (Add)	(None, 32, 32, 64)	0	['re_lu_4[0][0]', 'batch_normalization_6[0][0]']	
re_lu_6 (ReLU)	(None, 32, 32, 64)	0	['add_2[0][0]']	
conv2d_7 (Conv2D)	(None, 16, 16, 128)	73856	['re_lu_6[0][0]']	
batch_normalization_7 (BatchNo rmalization)	(None, 16, 16, 128)	512	['conv2d_7[0][0]']	
re_lu_7 (ReLU) [0][0]']	(None, 16, 16, 128)	0	['batch_normalization_7[0][0]']	
conv2d_9 (Conv2D)	(None, 16, 16, 128)	8320	['re_lu_6[0][0]']	
conv2d_8 (Conv2D)	(None, 16, 16, 128)	147584	['re_lu_7[0][0]']	
batch_normalization_9 (BatchNo rmalization)	(None, 16, 16, 128)	512	['conv2d_9[0][0]']	
batch_normalization_8 (BatchNo rmalization)	(None, 16, 16, 128)	512	['conv2d_8[0][0]']	
add_3 (Add) [0][0]', 8[0][0]']	(None, 16, 16, 128)	0	['batch_normalization_9[0][0]', 'batch_normalization_8[0][0]']	
re_lu_8 (ReLU)	(None, 16, 16, 128)	0	['add_3[0][0]']	
conv2d_10 (Conv2D)	(None, 16, 16, 128)	147584	['re_lu_8[0][0]']	
batch_normalization_10 (BatchN ormalization)	(None, 16, 16, 128)	512	['conv2d_10[0][0]']	
re_lu_9 (ReLU) 0[0][0]']	(None, 16, 16, 128)	0	['batch_normalization_10[0][0]']	

conv2d_11 (Conv2D)	(None, 16, 16, 128)	147584	['re_lu_9[0][0]']
batch_normalization_11 (Batch Normalization)	(None, 16, 16, 128)	512	['conv2d_11[0][0]']
add_4 (Add)	(None, 16, 16, 128)	0	['re_lu_8[0][0]', 'batch_normalization_11[0][0]']
re_lu_10 (ReLU)	(None, 16, 16, 128)	0	['add_4[0][0]']
conv2d_12 (Conv2D)	(None, 16, 16, 128)	147584	['re_lu_10[0][0]']
batch_normalization_12 (Batch Normalization)	(None, 16, 16, 128)	512	['conv2d_12[0][0]']
re_lu_11 (ReLU)	(None, 16, 16, 128)	0	['batch_normalization_12[0][0]']
conv2d_13 (Conv2D)	(None, 16, 16, 128)	147584	['re_lu_11[0][0]']
batch_normalization_13 (Batch Normalization)	(None, 16, 16, 128)	512	['conv2d_13[0][0]']
add_5 (Add)	(None, 16, 16, 128)	0	['re_lu_10[0][0]', 'batch_normalization_13[0][0]']
re_lu_12 (ReLU)	(None, 16, 16, 128)	0	['add_5[0][0]']
conv2d_14 (Conv2D)	(None, 8, 8, 256)	295168	['re_lu_12[0][0]']
batch_normalization_14 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_14[0][0]']
re_lu_13 (ReLU)	(None, 8, 8, 256)	0	['batch_normalization_14[0][0]']
conv2d_16 (Conv2D)	(None, 8, 8, 256)	33024	['re_lu_12[0][0]']

conv2d_15 (Conv2D)	(None, 8, 8, 256)	590080	['re_lu_13[0][0]']
batch_normalization_16 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_16[0][0]']
batch_normalization_15 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_15[0][0]']
add_6 (Add) 6[0][0]', 15[0][0]']	(None, 8, 8, 256)	0	['batch_normalization_16[0][0]', 'batch_normalization_15[0][0]']
re_lu_14 (ReLU)	(None, 8, 8, 256)	0	['add_6[0][0]']
conv2d_17 (Conv2D)	(None, 8, 8, 256)	590080	['re_lu_14[0][0]']
batch_normalization_17 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_17[0][0]']
re_lu_15 (ReLU) 7[0][0]']	(None, 8, 8, 256)	0	['batch_normalization_17[0][0]']
conv2d_18 (Conv2D)	(None, 8, 8, 256)	590080	['re_lu_15[0][0]']
batch_normalization_18 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_18[0][0]']
add_7 (Add) 18[0][0]']	(None, 8, 8, 256)	0	['re_lu_14[0][0]', 'batch_normalization_18[0][0]']
re_lu_16 (ReLU)	(None, 8, 8, 256)	0	['add_7[0][0]']
conv2d_19 (Conv2D)	(None, 8, 8, 256)	590080	['re_lu_16[0][0]']
batch_normalization_19 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_19[0][0]']
re_lu_17 (ReLU)	(None, 8, 8, 256)	0	['batch_normalization_19[0][0]']

9[0][0]']			
conv2d_20 (Conv2D)	(None, 8, 8, 256)	590080	['re_lu_17[0][0]']
batch_normalization_20 (Batch Normalization)	(None, 8, 8, 256)	1024	['conv2d_20[0][0]']
add_8 (Add)	(None, 8, 8, 256)	0	['re_lu_16[0][0]', 'batch_normalization_20[0][0]']
re_lu_18 (ReLU)	(None, 8, 8, 256)	0	['add_8[0][0]']
average_pooling2d (AveragePooling2D)	(None, 2, 2, 256)	0	['re_lu_18[0][0]']
flatten (Flatten)	(None, 1024)	0	['average_pooling2d[0][0]']
dense (Dense)	(None, 10)	10250	['flatten[0][0]']

=====

Total params: 4,344,842
Trainable params: 4,338,570
Non-trainable params: 6,272

None

The total parameters are 4.3M which are < 5M